

Phase 7 - OpenMP

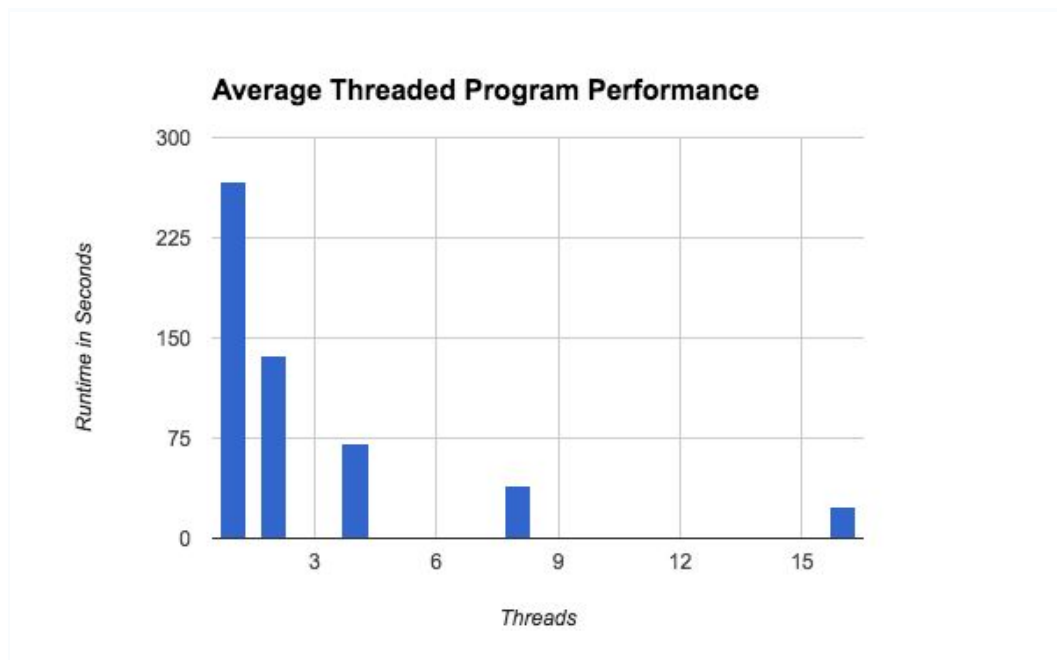
IT 515R - Scientific Computing

Tanner Satchwell

My approach in using OpenMP in my solver program was to implement it with as little change to my code as I could. This was a little difficult to do because, as we discussed in class, you don't spread out `#pragma omp` commands through functions. Generally best practice is to combine those functions so all of your `#pragma omp` code is in close proximity. Because I didn't want to change my code very much, I didn't follow best practice in this regard. This made it somewhat difficult in some instances to decide where I should put barriers to avoid deadlock and race conditions.

One change that I did end up making to my code was the way I checked to see if the grid was stable. Prior to implementing OpenMP, the `isStable` function returned as soon as the error was greater than epsilon. Because you can't break out of a for loop like that when you are using multiple threads, I had to move the error check outside of the for loop. This way, I was able to do a max reduction on the for loop, and then check if it was greater than epsilon afterward.

I was pleasantly surprised at how much of a performance increase using more threads added. Running my openmp program with a single thread took nearly five minutes, but adding more threads reduced that time by about half every time I doubled the thread count. See the following graph for a more detailed comparison.



Overall, I was pleased with the increased speed the program ran with more threads. If I were to do this project again however, I would probably do some major restructuring of my code to make threads work more intuitively. As it is, I'm still not 100% sure there are no race conditions. In Phase 8, I will most definitely restructure the code to make implementation of c++ threads easier.